

CUDA cheat sheet

Source code is in .cu files, which contain mixture of host (CPU) and device (GPU) code.

Declaring functions

<code>__global__</code>	declares kernel, which is called on host and executed on device
<code>__device__</code>	declares device function, which is called and executed on device
<code>__host__</code>	declares host function, which is called and executed on host

Declaring variables

<code>__device__</code>	declares device variable in global memory, accessible from all threads, with lifetime of application
<code>__constant__</code>	declares device variable in constant memory, accessible from all threads, with lifetime of application
<code>__shared__</code>	declares device variable in block's shared memory, accessible from all threads within a block, with lifetime of block

Vector types

`char1, uchar1, short1, ushort1, int1, uint1, long1, ulong1, float1`
`char2, uchar2, short2, ushort2, int2, uint2, long2, ulong2, float2`
`char3, uchar3, short3, ushort3, int3, uint3, long3, ulong3, float3`
`char4, uchar4, short4, ushort4, int4, uint4, long4, ulong4, float4`

`longlong1, ulonglong1, double1`
`longlong2, ulonglong2, double2`

`dim3`

Components are accessible as `variable.x`, `variable.y`, `variable.z`, `variable.w`.

Constructor is `make_<type>(x, ...)`, for example: `float2 xx = make_float2(1., 2.);`

`dim3` can take 1, 2, or 3 arguments:

`dim3 blocks1D(5);`
`dim3 blocks2D(5, 5);`
`dim3 blocks3D(5, 5, 5);`

Pre-defined variables

<code>dim3 gridDim</code>	dimensions of grid
<code>dim3 blockDim</code>	dimensions of block
<code>uint3 blockIdx</code>	block index within grid

uint3 threadIdx thread index within block

Kernel invocation

```
__global__ void kernel( ... ) { ... }
```

dim3 blocks(nx, ny, nz); // cuda 1.x has 1D and 2D grids, cuda 2.x adds 3D grids
dim3 threadsPerBlock(mx, my, mz); // cuda 1.x has 1D, 2D, and 3D blocks

```
kernel<<< blocks, threadsPerBlock >>>( ... );
```

Thread management

__threadfence_block(); wait until memory accesses are visible to block

__threadfence(); wait until memory accesses are visible to block and device

__threadfence_system(); wait until memory accesses are visible to block and device and host (2.x)

__syncthreads(); wait until all threads reach sync

Memory management

__device__ float* pointer;
cudaMalloc((void**) &pointer, size);
cudaFree(pointer);

// direction is one of cudaMemcpyHostToDevice or cudaMemcpyDeviceToHost
cudaMemcpy (dst_pointer, src_pointer, size, direction);
cudaMemcpyAsync(dst_pointer, src_pointer, size, direction, stream);

// using column-wise notation
// (the CUDA docs describe it for images; a “row” there equals a matrix column)
// _bytes indicates arguments that must be specified in bytes
cudaMemcpy2D (A_dst, lda_bytes, B_src, ldb_bytes, m_bytes, n, direction);
cudaMemcpy2DAsync(A_dst, lda_bytes, B_src, ldb_bytes, m_bytes, n, direction, stream);

// cublas makes copies easier for matrices, e.g., less use of sizeof
// copy x => y
cublasSetVector (n, elemSize, x_src_host, incx, y_dst_dev, incy);
cublasGetVector (n, elemSize, x_src_dev, incx, y_dst_host, incy);
cublasSetVectorAsync(n, elemSize, x_src_host, incx, y_dst_dev, incy, stream);
cublasGetVectorAsync(n, elemSize, x_src_dev, incx, y_dst_host, incy, stream);

// copy A => B
cublasSetMatrix (rows, cols, elemSize, A_src_host, lda, B_dst_dev, ldb);
cublasGetMatrix (rows, cols, elemSize, A_src_dev, lda, B_dst_host, ldb);
cublasSetMatrixAsync(rows, cols, elemSize, A_src_host, lda, B_dst_dev, ldb, stream);
cublasGetMatrixAsync(rows, cols, elemSize, A_src_dev, lda, B_dst_host, ldb, stream);

Also, malloc and free work inside a kernel (2.x), but memory allocated in a kernel must be deallocated in a kernel (not the host). It can be freed in a different kernel, though.

Atomic functions

```
old = atomicAdd ( &addr, value ); // old = *addr; *addr += value
old = atomicSub ( &addr, value ); // old = *addr; *addr -= value
old = atomicExch( &addr, value ); // old = *addr; *addr = value
```

```
old = atomicMin ( &addr, value ); // old = *addr; *addr = min( old, value )
old = atomicMax ( &addr, value ); // old = *addr; *addr = max( old, value )
```

```
// increment up to value, then reset to 0
// decrement down to 0, then reset to value
old = atomicInc ( &addr, value ); // old = *addr; *addr = ((old >= value) ? 0 : old+1 )
old = atomicDec ( &addr, value ); // old = *addr; *addr = ((old == 0) or (old > val) ? val : old-1 )
```

```
old = atomicAnd ( &addr, value ); // old = *addr; *addr &= value
old = atomicOr ( &addr, value ); // old = *addr; *addr |= value
old = atomicXor ( &addr, value ); // old = *addr; *addr ^= value
```

```
// compare-and-store
old = atomicCAS ( &addr, compare, value ); // old = *addr; *addr = ((old == compare) ? value : old)
```

Timer

```
wall clock cycle counter
clock_t clock();
```

cuBLAS

Matrices are column-major. Indices are 1-based; this affects result of `i<t>amax` and `i<t>amin`.
`#include <cublas_v2.h>`

```
cublasHandle_t handle;
cudaStream_t stream;

cublasCreate( &handle );
cublasDestroy( handle );
cublasGetVersion( handle, &version );
cublasSetStream( handle, stream );
cublasGetStream( handle, &stream );
cublasSetPointerMode( handle, mode );
cublasGetPointerMode( handle, &mode );
```

Compiler

nvcc, often found in `/usr/local/cuda/bin`

Flags common with cc

Short flag	Long flag	Output or Description
-c	--compile	.o object file
-E	--preprocess	on standard output

-M	--generate-dependencies	on standard output
-o <i>file</i>	--output-file <i>file</i>	
-I <i>directory</i>	--include-path <i>directory</i>	header search path
-L <i>directory</i>	--library-path <i>directory</i>	library search path
-l <i>lib</i>	--library <i>lib</i>	link with library
-lib		generate library
-shared		generate shared library
-pg	--profile	for gprof
-g <i>level</i>	--debug <i>level</i>	
-G	--device-debug	
-O <i>level</i>	--optimize <i>level</i>	

Some hardware constraints

	1.x	2.x
max x- or y-dimension of block	512	1024
max z-dimension of block	64	64
max threads per block	512	1024
warp size	32	32
max blocks per MP	8	8
max warps per MP	32	48
max threads per MP	1024	1536
max 32-bit registers per MP	16k	32k
max shared memory per MP	16 KB	48 KB
shared memory banks	16	32
local memory per thread	16 KB	512 KB
const memory	64 KB	64 KB

const cache	8 KB	8 KB
texture cache	8 KB	8 KB